

Joshua Zimdars

Ann Arbor, MI · 734-730-9150 · jzimdars@umich.edu

[linkedin.com/in/joshua-zimdars-1321a624b](https://www.linkedin.com/in/joshua-zimdars-1321a624b) · jzimdars-rtg.github.io/portfolio

I build and ship production AI systems — and then I make sure the people around me actually use them. For the past 18 months I've functioned as a forward-deployed engineer embedded inside a real customer environment: owning the full lifecycle from identifying broken workflows to writing the code, running evals, driving adoption, and staying until the new system is the default. My degree is in Sociology. Everything I know about AI engineering I built in production, under real constraints, with real users depending on the output. I'm not transitioning into this field — I'm already in it.

EXPERIENCE

Golden Limousine International — Ann Arbor, MI

Business Systems Specialist & Executive Assistant to the CEO and CFO

Nov 2024 – Present

Sole engineer at a 35-year-old, 100-market ground transportation SMB. Identify high-friction workflows, prototype rapidly, and ship production systems to internal stakeholders — functioning as a one-person forward-deployed engineering team inside a non-technical organization.

AI systems, agentic workflows, and production deployment

- Designed and shipped the **Quote Calculator** (*in active daily use by sales and dispatch*): a production agentic LLM workflow that parses inbound customer emails with **Gemini 2.0 Flash on Vertex AI**, routes across a **LangGraph multi-agent graph** (14 trip types × 14 vehicle classes), and outputs structured quotes — eliminating a previously manual, error-prone quoting process.
- Built and maintain `jz-genai-agent-toolkit` (github.com/jzimdars-RTG/JZ-genai-): an open-source vendor-agnostic agentic core in **JavaScript (LangGraph.js)** and **Python**; implements multi-provider LLM routing, ReAct, self-reflection, hierarchical delegation, RAG retrieval (Vertex AI text-embedding-004 + cosine similarity), and human-in-the-loop checkpointing.
- Instrumented a **Playwright-based LLM eval pipeline** across 12 parser scenarios after discovering the model was reliably wrong on multi-leg itineraries; re-prompting against eval output raised accuracy to production threshold — a full observability-and-correction loop on a live agentic system.
- Tracked **LLM-native metrics** per call: token consumption, cost-per-request, latency (p50/p95), and error rate via JSONL tracing; used data to drive prompt and routing decisions.
- Introduced AI-assisted workflows and prompt engineering practices across executives, sales, dispatch, and operations; led hands-on adoption sessions until teams preferred the new systems over legacy processes.

Engineering and platform

- Re-platformed the rate card into a **.NET 8 / EF Core / Azure SQL pricing API** (*9 controllers, 22 entities*) with a 10-phase product roadmap.
- Built a **real-time vehicle-tracking PWA** for U-M Hospital shuttles (*5,000+ monthly riders*) — cut customer status inquiries **40%** and freed **1,095 dispatch hours/year**.
- Built a **real-time vehicle detailing system** (*React 19 + Socket.io + Node*) with a 14-point QA checklist and tablet-optimized UI.
- Built a **Python/pandas passenger-analytics pipeline** now used in account renewal conversations with U-M Hospital.

- Shipped 8 internal applications collectively saving **2,800+ hours/year**.

Executive support and operations

- Prepare CEO/CFO weekly briefings and board-prep materials; lead cross-functional governance cadence across HR, IT, Safety, Fleet, and Operations.
- Cut IT support tickets **30%** by standardizing **60+ devices** under JAMF MDM; reduced deployment time **70%** with PowerShell CI/CD on Azure.
- Built a CDL driver recruitment site with full UTM/GA4 attribution in a day, giving leadership visibility into recruiting-channel performance across 4 job boards and 4 CDL roles.

University of Michigan, Institute for Social Research — Ann Arbor, MI

Research Assistant · May – Aug 2022

Synthesized international health-policy datasets into briefings for senior researchers; produced quantitative and qualitative analyses presented at academic forums.

Earlier

Server, Barton Hills Country Club (2024) and Palio Restaurant (2021–2023) — consultative, top-of-floor sales performer; built repeat clientele in high-touch hospitality settings.

Volunteer, Delonis Homeless Shelter (2022).

EDUCATION

University of Michigan, Ann Arbor, MI

B.A., Sociology of Health and Medicine · May 2024 · Michigan Club Lacrosse (2 yrs)

Equivalent practical experience in AI/ML systems engineering, agentic workflow development, and cloud platform deployment — acquired entirely through production work.

SELECTED PROJECTS · **JZIMDARS-RTG.GITHUB.IO/PORTFOLIO**

- **GenAI Agent Toolkit** — vendor-agnostic LangGraph agentic core in **JavaScript and Python**; multi-provider LLM routing, self-reflection, **RAG retrieval** (Vertex AI embeddings + cosine similarity), ReAct, hierarchical delegation, human-in-the-loop. github.com/jzimdars-RTG/JZ-genai-
- **Quote Calculator** (*in active production use*) — AI-assisted quoting and dispatch platform; Gemini 2.0 Flash on Vertex AI, LangGraph multi-agent graph, 14 trip types × 14 vehicles, Playwright eval suite, JSONL observability.
- **Pricing API** — .NET 8 / EF Core / Azure SQL; 9 controllers, 22 entities, 10-phase roadmap.
- **Shuttle Tracking PWA** — real-time vehicle map for 5K+ monthly U-M Hospital riders; reduced inquiry volume 40%.
- **Driver Recruitment Site** — UTM/GA4 attribution stack. Live at charterbusdrivers.com.

SKILLS

AI engineering & agentic systems — prompt engineering, Retrieval-Augmented Generation (RAG), multi-agent orchestration (ReAct, self-reflection, hierarchical delegation), LLM evaluation pipelines, human-in-the-loop, MCP patterns

Languages — JavaScript (Node.js, React 19), **Python** (LangGraph, pandas, data pipelines), C# / .NET 8, PowerShell, SQL, VBA

AI platforms — Google Vertex AI (Gemini 2.0 Flash, text-embedding-004), Azure AI Inference (Kimi K2), OpenAI, Claude, GitHub Copilot

Agentic frameworks — LangGraph.js, LangGraph (Python), langchain-google-vertexai; patterns: ReAct, self-reflection, human-in-the-loop

LLM-native observability — per-call token accounting, cost-per-request, latency (p50/p95), JSONL tracing, error rate

Cloud platforms — Google Cloud Platform (Vertex AI, GCP), Azure (Functions, Static Web Apps, SQL, Maps, AI Inference, MDM/Intune)

Engineering — REST APIs, Socket.io, ES modules, .NET 8 / EF Core, PowerShell CI/CD, Playwright (LLM eval), JAMF Now MDM

Data & analytics — SQL, Python/pandas, Excel modeling, GA4, UTM attribution, dashboarding

Strategy & operations — pricing strategy, unit economics, stakeholder alignment, change management, governance cadence

Certifications — MaintainX Administrator Foundations (2025) · Deloitte Future of Work Institute (2023)